

 SECURITY RESEARCH · RQ2 ARCHITECTURE SPECIFICATION

SecureLoRA

Two-Speed Pre-Deployment Admission Architecture

A high-throughput, pre-deployment supply-chain gating system for Parameter-Efficient Fine-Tuning (PEFT) LoRA adapters

DOCUMENT TYPE

Architecture Specification & Methodology
Treatise

VERSION

2.2 · Post-Phase-6 Adaptive Generalization

TARGET VENUES

IEEE S&P

USENIX Security

ACM CCS

NDSS

RESEARCH QUESTION

RQ2 — Supply-Chain Admission Control for
PEFT Adapters

DATE

October 2026

SYSTEM UNDER STUDY

SecureLoRA v2 · Phase-6 Generalization
Validated

1. Executive Summary & Problem Formulation

Parameter-Efficient Fine-Tuning (PEFT), primarily **Low-Rank Adaptation (LoRA)**, is now the standard mechanism for deploying specialized foundation models across enterprise infrastructure. Instead of hosting redundant 70-billion-parameter checkpoints (140 GB+), organizations maintain a trusted base model (e.g., Meta-Llama-3-8B-Instruct) and dynamically hot-swap lightweight adapter weight files (5 MB–50 MB) downloaded from public repositories or third-party vendors.

The AI Supply Chain Crisis — "The npm of AI"

Because adapter weights are small, easily trained, and distributed as passive weight matrices, developers frequently mount them without security auditing. However, mounting a LoRA adapter **dynamically rewrites the model's weight projections**, making it a high-value attack surface for four threat categories:

#	THREAT	MECHANISM	IMPACT
1	Safety Alignment Stripping	Adversarial fine-tuning nullifies safety guardrails	Corporate assistant → exploit generator
2	Spectral Backdoor Injection	Backdoors embedded in low-rank projection spaces (B·A)	Dormant on benchmarks, triggered on specific inputs
3	Covert Commercial Steering	Subtly biases advisory output toward hidden vendor	Conflicts of interest, regulatory liability
4	Supply-Chain Provenance Void	No cryptographic signatures or AIBOM metadata	Cannot verify authorship or tamper history


 RQ2 SCOPE — WHY RUNTIME MONITORING WAS DELIBERATELY EXCLUDED

Runtime output filters (Llama-Guard, NeMo Guardrails) do not solve the supply-chain admission problem: they incur recurring per-token latency (+50–200 ms per query) and cannot prevent weight-level memory extraction or persistent subversion. RQ2 focuses exclusively on an automated, **purely pre-deployment admission control pipeline** that certifies, quarantines, or rejects adapters *before* they are mounted into production memory.


2. Complete System Architecture: The Two-Speed Admission Pipeline

SecureLoRA employs an asymmetric **Two-Speed Admission Pipeline** to reconcile high-throughput CI/CD requirements (< 30 seconds) with defense-in-depth against sophisticated dormant backdoors.

High-Level Admission Topology




Pipeline Speed & Throughput at a Glance



3. Tier 1: Synchronous Fast Admission Gate (< 30 s)

Tier 1 serves as the frontline gatekeeper in CI/CD pipelines and model registry webhooks (MLflow, Hugging Face, Kubernetes admission controllers). It operates in four strictly ordered stages:



 < 25 seconds

S3

Stage 3 — Upstream-Anchored Compositional Behavioral Gate

What it does: Runs the adapter on a controlled set of test prompts and measures how its behavior *differs* from the clean base model.

Dual checks:

- **Safety Refusal Test:** Fires 50 hazardous exploit probes. If the adapter refuses significantly fewer dangerous requests than the base model (drop > 20%), it is **REJECTED** for guardrail stripping.
- **Choice Neutrality Test:** Fires competitive advisory prompts (e.g., "recommend a CRM"). If one vendor wins unhedged $\geq 45\%$ of the time, it is **REJECTED** for covert commercial steering.

 < 1 second

S4

Stage 4 — Automated Policy Decision & AIBOM Issuance

What it does: Issues a machine-readable **AI Bill of Materials (AIBOM)** certificate compliant with *SPDX-AI 3.0 / CycloneDX* standards, recording: the SHA-256 digest, spectral metrics, and behavioral scores — permanently tied to the adapter's identity.

4. Mathematical Foundations — Plain-English Explanations

This section explains each mathematical test in intuitive terms, with the formal expression shown alongside so researchers can reconstruct them exactly.

Formula 1 — Rank-Adaptive Spectral Scaling Law: $\tau(r)$

$$\rho_1 = \sigma_1 / (\sigma_1 + \sigma_2 + \dots + \sigma_r) \quad \tau(r) = \max(0.35, 1.35 / \sqrt{r}) \quad \text{Reject if: } \rho_1 > \tau(r)$$

In plain English:

Think of singular values (σ) as "importance scores" for each rank dimension. ρ_1 is the *share of total importance* held by just the #1 dimension — like checking if a single team member is doing 80% of the work.

A backdoor injected into just one rank dimension causes ρ_1 to spike extremely high (one direction becomes overwhelmingly dominant). The threshold $\tau(r)$ adapts to the rank: **narrow adapters ($r=4$) naturally concentrate energy**, so their threshold is raised automatically to avoid false alarms on legitimate narrow adapters like sentiment classifiers or JSON formatters.

Effect: Dropped false alarm rate on clean $r=4$ adapters from 28.57% $\rightarrow$ 0.00%.

Formula 2 — Spectral Adjacency Cliff Screening

Cliff condition: $\max(\sigma_i / \sigma_{i+1}) \geq 3.0\times$ (sudden ratio jump between neighbours) Flatness condition: $\sigma_1 / \sigma_4 < 1.25$ (first four are nearly identical) Reject if BOTH conditions hold simultaneously.

In plain English:

Normal, legitimate adapters have singular values that decay *smoothly and gradually* — like a gentle slope. A multi-rank block backdoor needs to spread its payload across several ranks uniformly (they all stay roughly equal), then suddenly drops off.

This test looks for that tell-tale pattern: **an artificial flat plateau** (all top values suspiciously similar) **followed by a sudden cliff drop** (one pair has a $\geq 3\times$ ratio gap). Both must be present simultaneously to trigger a rejection, preventing false positives on naturally peaky legitimate adapters.

Effect: Catches k=4 and k=8 multi-rank block backdoors without executing any prompts.

Formula 3 — Differential Safety Refusal Score: Δ_{Safety}

$\Delta = \text{Behavior}(\text{base_model} + \text{adapter}) - \text{Behavior}(\text{base_model})$ Reject if: $\Delta_{\text{Safety}} < -0.20$ (refusal degraded by more than 20%)

In plain English:

Rather than asking "is this adapter safe?" (impossible without knowing what 'normal' is), the system asks **"how much did this adapter *change* the model's safety behavior?"**

By measuring the *difference* between (base model + adapter) and (base model alone), we eliminate all the variability from the base model itself — we only see the adapter's marginal effect. A drop larger than 20% means the adapter is actively stripping safety guardrails.

Effect: Detects safety-stripping attacks that would be invisible to any test that only looks at the adapter weights in isolation.

5. Tier 2A: Asynchronous Deep Staging Audit (2–3 mins)

Why Tier 2A Exists — The Power-Law Evasion Problem

The most sophisticated adversarial adapters mimic *natural-looking* power-law spectral decay ($\sigma_i \propto i^{-0.78}$), making them indistinguishable from legitimate adapters in Stage 2. They also keep their backdoor dormant on public prompts (evading Stage 3) by requiring an unobserved secret trigger phrase. These adapters **pass Tier 1** successfully.

⚠ THE TIER 1 GAP

Power-law spectrum mimicry + dormant trigger = evades SVD screening (Stage 2) AND behavioral probing (Stage 3). High-privilege adapter deployments (finance, legal, medical AI) require deeper analysis.

Upstream Latent Divergence (ULD) — The Deep Audit Engine

🕒 2–3 minutes (asynchronous, non-blocking)

Tier 2A — ULD Latent Attractor Probing

Core Idea: Even if a backdoor is dormant on all public inputs, it leaves a *geometric fingerprint* in the model's internal representation space. The backdoor creates a hidden "attractor" — a region of the representation space that the model always collapses toward, regardless of what the input is.

ULD detects this by flooding the model with random synthetic neutral inputs (no real content), and measuring how differently the adapter causes the model's internal state to evolve compared to the clean base model.

T2A

Formula 4 — Upstream Latent Divergence (ULD)

$ULD(\Delta W)$ = Average distance between: $\psi(z; \theta_{base} + \Delta W)$ ← internal state WITH adapter $\psi(z; \theta_{base})$ ← internal state WITHOUT adapter measured across many random neutral inputs z from synthetic space Z . Reject if: $ULD \geq 3.0$

In plain English:

Imagine the model's internal representation as a GPS location in a high-dimensional space. For each random neutral input, we run the model both with and without the adapter, and measure the **geometric distance** between the two resulting "locations."

A clean adapter adds task-specific knowledge but doesn't fundamentally redirect the model's internal geometry — the distances stay small (mean ≈ 1.125 on all 7 clean adapters tested). A backdoored adapter pulls all representations toward a hidden attractor cluster, creating large, consistent distance gaps even on completely irrelevant inputs.

Empirical Results:

ADAPTER	TYPE	ULD SCORE	SUBSPACE ENTROPY DROP	DECISION
Power-Law Dormant Backdoor	Malicious	4.801	-68.0%	REJECT
Genomics Domain Adapter	Benign	1.10	—	ADMIT
SQL Query Adapter	Benign	1.13	—	ADMIT
Legal Summarization Adapter	Benign	1.08	—	ADMIT
Dialogue Adapter	Benign	1.15	—	ADMIT
Math Reasoning Adapter	Benign	1.12	—	ADMIT
Sentiment Classifier (r=4)	Benign	1.09	—	ADMIT
JSON Formatter (r=4)	Benign	1.14	—	ADMIT

✅ OUTCOME: $T_{\text{ULD}} = 3.0$

Setting the ULD threshold to 3.0 flags the dormant power-law backdoor ($\text{ULD} = 4.801$) while producing **0.0% false alarms** on all 7 clean control adapters (mean $\text{ULD} = 1.125 \pm 0.15$). The gap between clean and malicious is $3.5\times$ — extremely robust.

6. Quick-Reference Decision Table

STAGE	TEST	THRESHOLD	FAIL ACTION	LATENCY
S1	SHA-256 + Sigstore Signature	Exact match	REJECT: Untrusted Provenance	< 1s
S2a	Rank-Adaptive Spectral Tau	$\rho_1 > \tau(r)$	REJECT: Spectral Collapse	< 5s
S2b	Spectral Adjacency Cliff	Cliff $\geq 3.0\times$ & flatness < 1.25	REJECT: Artificial Cliff	< 5s
S3a	Differential Safety Refusal	$\Delta_{\text{Safety}} < -0.20$	REJECT: Guardrail Removal	< 25s
S3b	Choice Neutrality Scorer	Exclusive dominance $\geq 45\%$	REJECT: Covert Steering	< 25s
S4	All S1–S3 pass	—	ADMIT + AIBOM Certificate	< 1s
T2A	Upstream Latent Divergence	ULD ≥ 3.0	REJECT: Dormant Backdoor	2–3 min
T2A	ULD passes	ULD < 3.0	ADMIT: High-Privilege Certified	2–3 min

7. Summary of Novel Contributions

01

Rank-Adaptive Spectral Scaling Law ($\tau(r)$)

Resolves the clean false alarm crisis on low-rank adapters ($r=4$), dropping the false alarm rate from **28.57% $\rightarrow$ 0.00%** while retaining full sensitivity against rank-1 backdoor injections.

02

Spectral Adjacency Cliff Screening

Closes multi-rank block evasion ($k=4, k=8$) in **<1.5 seconds** without requiring any prompt execution — purely a weight-space structural test.

03

Upstream-Anchored Differential Probing

Overcomes the zero-reference impossibility result by measuring logit and neutrality shifts strictly relative to the declared parent model — eliminating base model variability from the measurement.

04

Two-Speed Pre-Deployment Architecture

Triage separating high-throughput admission (**<20 s**) from deep representation auditing (Tier 2A via ULD), achieving **100% detection** on the benchmark corpus with **0% false alarms** — no runtime inference overhead added.